

Customer Contract Management Portal — System Architecture

Diagram

```
graph TD
    subgraph Frontend
        A[React SPA]
    end

    subgraph Backend
        B[API Gateway / Reverse Proxy]
        C[Backend App (FastAPI / Express)]
        CA[Action Service]
        CE[Event Ingest Service]
        AU[Auth / SSO (SAML/OIDC)]
        NOTIF[Notification Service]
        ANALYTICS[Optional AI / Analytics]
    end

    subgraph DataLayer
        DB[SQLite (Primary) \n customers, contracts, events, notes]
        AUDIT[Immutable Audit Store / Append-only Log]
        SEARCH[Optional Full-text Search / Index]
        DW[Optional Data Warehouse]
    end

    subgraph Observability
        LOGS[Centralized Logs]
        TRACES[Distributed Tracing]
        METRICS[Metrics & Alerts]
    end

    A -->|HTTPS REST| B
    B --> C
    C -->|SQL| DB
    C -->|write| AUDIT
    CE -->|ingest events (HTTP/gRPC)| C
    CE -->|append| DB
    C --> SEARCH
    NOTIF -->|email/IM| ExternalSystems[(Email / Internal Chat)]
    C --> AU
```

```

ANALYTICS -.-> DB
ANALYTICS -.-> C

C --> LOGS
C --> TRACES
C --> METRICS
CE --> LOGS
B -->|TLS termination| AU

%% Endpoints
classDef endpoints fill:#f9f,stroke:#333;
subgraph Endpoints[Main REST Endpoints]
    E1[/customers]
    E2[/contracts]
    E3[/contracts/{id}]
    E4[/events]
    E5[/actions]
end
B --> E1
B --> E2
B --> E3
B --> E4
B --> E5

%% Future integrations
ExternalAI[(External AI / Risk Scoring)]
ExternalIDP[(Enterprise IdP)]
ExternalEventSources[(Event Producers / Kafka)]
ExternalSIEM[(SIEM / Security)]

CE --> ExternalEventSources
ANALYTICS --> ExternalAI
LOGS --> ExternalSIEM
AU --> ExternalIDP

```

Textual explanation

High-level overview

This system is a classic three-tier web app:

- **Frontend:** React single-page application providing pages (Dashboard, Customers, Contracts, Contract Details, Event Log). Communicates with backend via HTTPS REST API; uses token-based session from SSO (OIDC/SAML).

- **Backend:** Monolithic or modular service implemented in **FastAPI (Python)** or **Express (Node.js)** exposing the main REST endpoints. Implements business logic (state machine for contract lifecycle), RBAC checks, input validation, audit recording, and event ingestion.
- **Database:** SQLite as requested for Phase 1. Stores `customers`, `contracts`, `events`, and `notes`. An append-only **Audit Store** is used to keep tamper-evident records for approvals, flags, and comments.

Components and responsibilities

- 1. React SPA (Frontend)**
2. Pages: Dashboard, Customers, Contracts, Contract Details, Event Log.
3. Uses OIDC/SAML SSO for authentication and receives a scoped access token and user role claims.
4. Calls backend REST endpoints, displays data tables and charts, enforces client-side input validation and masking rules (partial national ID, emails).
5. Handles optimistic UI updates for comments/flags and displays action results with proper error handling.
- 6. API Gateway / Reverse Proxy**
7. TLS termination, basic rate limiting, routing to backend instances, and static asset hosting.
8. Enforces authentication early, validates tokens with IdP, and injects user context (user id, roles) into requests.
- 9. Backend App (FastAPI / Express)**
10. Exposes REST endpoints: `/customers`, `/contracts`, `/contracts/{id}`, `/events`, `/actions`.
11. Responsibilities: input validation, authorization (RBAC + row-level checks), business workflows (contract state machine), comments/flags/actions handling, CSV export generation, and pagination.
12. Persists all mutating operations to the **Audit Store** in append-only fashion and writes normalized data to SQLite.
13. Event ingestion endpoint(s) accept external events and persist into `events` table after schema validation.
- 14. Action Service (module / queue consumer)**
15. Processes long-running or side-effectful actions (e.g., notifications, heavy exports, or integration calls). Keeps API responsive by returning accepted responses and doing work asynchronously.
- 16. Event Ingest Service**
17. Dedicated endpoint for event producers. Validates schema, deduplicates, and writes into `events` table. If event volume grows, this can buffer into Kafka or a queue.
- 18. Auth / SSO (Enterprise IdP)**

19. SAML/OIDC provider for single sign-on and MFA. Backend validates tokens and applies role-based claims.

20. **Notification Service**

21. Sends email or internal notifications when an approver is needed or when @mentions occur.

22. **Analytics / AI (Optional)**

23. Offline or microservice that consumes event/contract data to produce risk scores, trends, and suggested actions. Integrates with backend via an authenticated API or database exports.

24. **Data Layer**

25. **SQLite** as primary Phase-1 datastore. Tables: `customers`, `contracts`, `events`, `notes`. Use WAL mode and file-level backups for durability.

26. **Audit Store**: could be a separate append-only file, secure blob store, or a write-once S3 bucket with integrity hashes to meet tamper-evidence and retention.

27. **Optional Search Index** (e.g., Tantivy / Elasticsearch in future) for fast full-text search over contracts and comments.

28. **Observability**

- Centralized logs, metrics, and tracing (e.g., Prometheus + Grafana, Jaeger). Export logs to SIEM for compliance.

Data flow (typical scenarios)

Contract approval flow 1. Approver loads Contract Details page in React. Frontend calls `GET /contracts/{id}`. 2. Backend authorizes the call, reads from SQLite, returns contract data and comments. 3. Approver clicks Approve → frontend `POST /actions` with `{action: 'approve', contract_id, comment}`. 4. Backend validates role, records `ContractAction` in Audit Store and `contracts` table status change, then returns success. Backend enqueues notifications for interested parties. 5. Notification Service sends internal message or email.

Event ingestion flow 1. External system posts event to `/events` (or to Event Ingest Service). 2. Event service validates schema and writes event to `events` table and `audit` if required. 3. Dashboard or Event Log queries `GET /events` with filters to display recent events.

REST endpoints (concise mapping)

- `GET /customers` — list, search, filters, pagination
- `GET /customers/{id}` — detail with associated contracts
- `GET /contracts` — list and filters
- `GET /contracts/{id}` — contract details, comments, flags, history
- `GET /events` — query events by customer, type, time range

- POST /actions — {action: 'approve' | 'reject' | 'flag' | 'note', payload: {...}}
(authz required)
- POST /events — event ingestion (secure, authenticated producer)
- GET /export — CSV export endpoint (admin/analyst permissions)

Best practices — Security

- Use enterprise SSO (SAML/OIDC) + MFA. Validate tokens at the gateway.
- Enforce RBAC and row-level authorization in the backend for all mutating/read operations.
- Encrypt transport with TLS 1.2+ and enable HSTS.
- Encrypt sensitive fields at rest (AES-256) — especially National ID and any PII; use field-level encryption if possible.
- Store secrets in a vault and rotate regularly.
- Implement CSRF protections and strong Content Security Policy in frontend.
- Input validation and strict schema for events; sanitize and scan attachments for malware.
- Immutable audit logs with write-once semantics; sign exported CSVs using a hash for integrity.

Best practices — Scalability & Availability

- SQLite is fine for Phase 1 and low concurrency proofs of concept. For production scale targets (1M customers, 500M events) plan migration to a client-server RDBMS (Postgres, cloud managed RDS) and use partitioning for events.
- Introduce read replicas and a caching layer (Redis) for hot reads (dashboard counts, frequent queries).
- Move event ingestion to an append-only distributed log (Kafka or cloud streaming) to decouple producers from the portal.
- Use horizontal scaling for backend (stateless containers) behind the API gateway.
- Use async workers (Celery / RQ / Bull) for heavy tasks (exports, notifications).

Best practices — Modularity & Maintainability

- Keep bounded contexts: split services (ingest, actions, analytics) as separate modules/services.
- Define clear API contracts and use OpenAPI / AsyncAPI for documentation and client generation.
- Write comprehensive tests: unit, integration, and contract tests for the event schema.
- Feature flags for gating dashboard widgets and edit windows.
- CI/CD with automated security scans (SAST/DAST) and deployment pipelines per environment.

Deployment considerations

- Environments: dev, staging, prod. Use IaC (Terraform) for infra.
- Back up SQLite files regularly; ensure consistent snapshots when writing (use WAL + file system snapshot) or migrate to managed DB for safer backups.
- Retention policy: keep audit and events per compliance (7 years); consider cold storage for older events.

Notes & trade-offs

- **SQLite advantages:** simple to deploy, low ops overhead for Phase 1. **Limitations:** concurrency, scalability, backups and multi-tenant safety — plan to migrate to Postgres or cloud RDS for scale and reliability.
 - **Audit requirements:** implement append-only audit store separate from writable DB to meet tamper-evident requirements.
 - **Performance:** heavy event volumes require sharding/partitioning — implement time-based partitions and/or move events to a purpose-built event store.
-

If you'd like, I can: - Produce a sequence diagram for the contract approval flow. - Produce an infrastructure diagram (cloud provider specific — AWS/Azure/GCP) with recommended managed services. - Generate Kubernetes manifests / Dockerfile + sample FastAPI or Express skeleton implementing the endpoints.

Tell me which of those you'd like next and I'll add it directly into this canvas.